

End of course summative assessment:

# **Machine Learning Approaches for Breast Cancer Detection: A Comparative Study**

**By Kai Bowers**

*Student Number: 230319794*

*Course: CSM010-2023-APR*

***2086 Words***

# **Table of Contents**

1. Introduction
  1. Research motivation
  2. Local setup
2. The data
  1. Preparing the dataset
  2. Exploring the dataset
3. Building ML algorithms
  1. Supervised Model
  2. Un-supervised Model
  3. Semi-supervised Model
4. Constructing and Selecting Features
5. Evaluating models and analysing the results
6. Conclusion
7. Sources

## **Chapter 1.1 : Introduction : Research motivation**

With approximately 72,000 annual cases, breast cancer is by far the most common cancer in women. [1]

It is complicated and can be caused by many factors, like genes, environment and lifestyle. Machine learning can help with studying breast cancer, as it can analyze patterns and connections that humans might miss. AI can help us learn more about the disease, diagnose it accurately, and make better treatment decisions.

One important thing in breast cancer is finding it early. When we find breast cancer early, it's easier to treat and can save lives. AI and machine learning can look at pictures of the breast, like mammograms or ultrasounds, and find signs that might mean cancer is there.

AI and machine learning are also helpful in predicting what might happen with the disease. Doctors want to know how fast the cancer might grow or if it will come back after treatment. Using AI, scientists can analyze a lot of information, like genes and clinical data, to make predictions. This helps doctors make better decisions about treatment and tell patients what to expect.

In this paper, I am going to explore the dataset given by UI Irvine titled "Breast Cancer Wisconsin, advanced models for cancer detection". I will focus on utilizing the Breast Cancer Wisconsin Diagnostic dataset [2] to train multiple machine learning models capable of detecting breast cancer. The dataset provides us with 30 features which were computed from three digitized images of a fine needle aspirate (FNA) of a breast mass. The classification of each data point (row) will be a diagnoses of either M (malignant ) or B (benign). There are 569 rows in the dataset.

The ten features of a single image are : radius , texture, perimeter , area , smoothness , compactness , concavity , concave points , symmetry, fractal dimension. We are given all ten features for each of the three images. There is also a row for a ID and the diagnoses.

## **Chapter 1.2 : Introduction : Local setup**

I created all code locally using VS Code, Python 3 and Windows 11.

To run, simply clone the repo, install python and the required PIP packages.

All following screenshots of the jupyter notebooks were created by me and are visible in the repository. [3]

## Chapter 2.1 : The data : Preparing the dataset

The raw data is given to us in the “wdbc.data”-file. In the jupyter notebook “Data pre-processing” I prepared the data for my machine learning algorithms. First, I created a copy of the file and added the feature labels as the first row of the CSV dataset. The features are available at the source of the dataset. [2]

```
# The goal of this Notebook is the process the wdbc.data into a wdbc.csv

# First, Import Libraries

import shutil
import pandas as pd

# Copy The Original Raw Dataset to "wdbc.csv"

shutil.copyfile("wdbc.data", "wdbc.csv")

'wdbc.csv'

# These Features are defined at https://archive.ics.uci.edu/dataset/17/breast+cancer+wisconsin+diagnostic

Features = "ID,Diagnosis,radius1,texture1,perimeter1,area1,smoothness1,compactness1,concavity1,concave_points1,symmetry1,fractal_
4

# Add The Features as the first row of the CSV

with open('wdbc.csv', 'r') as original: data = original.read()
with open('wdbc.csv', 'w') as modified: modified.write(Features + "\n" + data)
```

I then removed the ID collum and normalized the data. There was no need for the ID field as it didn't provide additional information to the models. The source assured me that all values were filled and none were missing, which I double checked and confirmed to be correct.

```
In [10]: # Remove the ID Collum

In [11]: df = pd.read_csv('wdbc.csv')
first_column = df.columns[0]
df = df.drop([first_column], axis=1)
df.to_csv('wdbc.csv', index=False)

In [12]: # Normalize each collum

In [13]: import pandas as pd
from sklearn.preprocessing import MinMaxScaler

data = pd.read_csv('wdbc.csv')
cols_to_normalize = data.columns[1:]
scaler = MinMaxScaler()
data[cols_to_normalize] = scaler.fit_transform(data[cols_to_normalize])
data.to_csv('wdbc.csv', index=False)

In [ ]: # Check if every value in the csv is set

In [15]: df = pd.read_csv("wdbc.csv")
is_empty = df.isnull().values.any()
if is_empty:
    print("There are missing values in the CSV file.")
else:
    print("All values in the CSV file are set.")

All values in the CSV file are set.
```

Running this Jupyter Notebook creates the wdbc.csv which will be used in the next tasks.

## Chapter 2.2 : The data : Exploring the dataset

Before starting the machine learning process, I wanted to understand the dataset. To do this I created the “Data understanding.ipynb” notebook in which I ran multiple tests that would help me decide which type of learning algorithms to use for my task.

```
# Imports
```

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns #Install
from scipy.stats import chi2_contingency, ttest_ind
```

```
# Print the Head of the CSV
```

```
data = pd.read_csv('wdbc.csv')
print(data.head())
```

	Diagnosis	radius1	texture1	perimeter1	area1	smoothness1	\
0	M	0.521037	0.022658	0.545989	0.363733	0.593753	
1	M	0.643144	0.272574	0.615783	0.501591	0.289880	
2	M	0.601496	0.390260	0.595743	0.449417	0.514309	
3	M	0.210090	0.360839	0.233501	0.102906	0.811321	
4	M	0.629893	0.156578	0.630986	0.489290	0.430351	

	compactness1	concavity1	concave_points1	symmetry1	...	radius3	\
0	0.792037	0.703140	0.731113	0.686364	...	0.620776	
1	0.181768	0.203608	0.348757	0.379798	...	0.606901	
2	0.431017	0.462512	0.635686	0.509596	...	0.556386	
3	0.811361	0.565604	0.522863	0.776263	...	0.248310	
4	0.347893	0.463918	0.518390	0.378283	...	0.519744	

	texture3	perimeter3	area3	smoothness3	compactness3	concavity3	\
0	0.141525	0.668310	0.450698	0.601136	0.619292	0.568610	
1	0.303571	0.539818	0.435214	0.347553	0.154563	0.192971	
2	0.360075	0.508442	0.374508	0.483590	0.385375	0.359744	
3	0.385928	0.241347	0.094008	0.915472	0.814012	0.548642	
4	0.123934	0.506948	0.341575	0.437364	0.172415	0.319489	

	concave_points3	symmetry3	fractal_dimension3
0	0.912027	0.598462	0.418864
1	0.639175	0.233590	0.222878
2	0.835052	0.403706	0.213433
3	0.884880	1.000000	0.773711
4	0.558419	0.157500	0.142595

[5 rows x 31 columns]

Calling `print(data.head())` allows us to view the first five lines of the normalized data.

```
# Summary statistics of the data
```

```
print(data.describe())
```

	radius1	texture1	perimeter1	area1	smoothness1	\
count	569.000000	569.000000	569.000000	569.000000	569.000000	
mean	0.338222	0.323965	0.332935	0.216920	0.394785	
std	0.166787	0.145453	0.167915	0.149274	0.126967	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.223342	0.218465	0.216847	0.117413	0.304595	
50%	0.302381	0.308759	0.293345	0.172895	0.390358	
75%	0.416442	0.408860	0.416765	0.271135	0.475490	
max	1.000000	1.000000	1.000000	1.000000	1.000000	


	compactness1	concavity1	concave_points1	symmetry1	\
count	569.000000	569.000000	569.000000	569.000000	
mean	0.260601	0.208058	0.243137	0.379605	
std	0.161992	0.186785	0.192857	0.138456	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.139685	0.069260	0.100944	0.282323	
50%	0.224679	0.144189	0.166501	0.369697	
75%	0.340531	0.306232	0.367793	0.453030	
max	1.000000	1.000000	1.000000	1.000000	

	fractal_dimension1	...	radius3	texture3	perimeter3	\
count	569.000000	...	569.000000	569.000000	569.000000	
mean	0.270379	...	0.296663	0.363998	0.283138	
std	0.148702	...	0.171940	0.163813	0.167352	
min	0.000000	...	0.000000	0.000000	0.000000	
25%	0.163016	...	0.180719	0.241471	0.167837	
50%	0.243892	...	0.250445	0.356876	0.235320	
75%	0.340354	...	0.386339	0.471748	0.373475	
max	1.000000	...	1.000000	1.000000	1.000000	

```
In [5]: # Visualize the target distribution
```

```
data['Diagnosis'].value_counts().plot(kind='bar')
plt.title('Target Distribution')
plt.xlabel('Target')
plt.ylabel('Count')
plt.show()
```



It is important to note that the amount of targets resulting in "B" was almost double the amount of targets resulting in "M". This made me want to analyze both outputs using t-tests and chi-square tests.

```
In [9]: # t-tests
numerical_cols = data.select_dtypes(include='number').columns
for col in numerical_cols:
    target_M = data[data['Diagnosis'] == 'M'][col]
    target_B = data[data['Diagnosis'] == 'B'][col]
    t_stat, p_value = ttest_ind(target_M, target_B)
    print(f'T-test for column {col}:')
    print(f'T-statistic: {t_stat}')
    print(f'P-value: {p_value}')
    print('---')
```

```
T-test for column radius1:
T-statistic: 25.43582161005704
P-value: 8.465940572264348e-96
---
T-test for column texture1:
T-statistic: 10.867201081464334
P-value: 4.0586360478983358e-25
---
T-test for column perimeter1:
T-statistic: 26.405212979192687
P-value: 8.436251036172328e-101
---
T-test for column area1:
T-statistic: 23.93868723569098
P-value: 4.7345643103078834e-88
---
T-test for column smoothness1:
T-statistic: 9.14609880814903
P-value: 1.0518503592032693e-18
```

```
In [10]: # chi-square test
categorical_cols = data.select_dtypes(include='object').columns
for col in categorical_cols:
    contingency_table = pd.crosstab(data[col], data['Diagnosis'])
    chi2, p_value, _, _ = chi2_contingency(contingency_table)
    print(f'Chi-square test for column {col}:')
    print(f'Chi2 statistic: {chi2}')
    print(f'P-value: {p_value}')
    print('---')
```

```
Chi-square test for column Diagnosis:
Chi2 statistic: 564.7302404341926
P-value: 7.86394182828703e-125
---
```

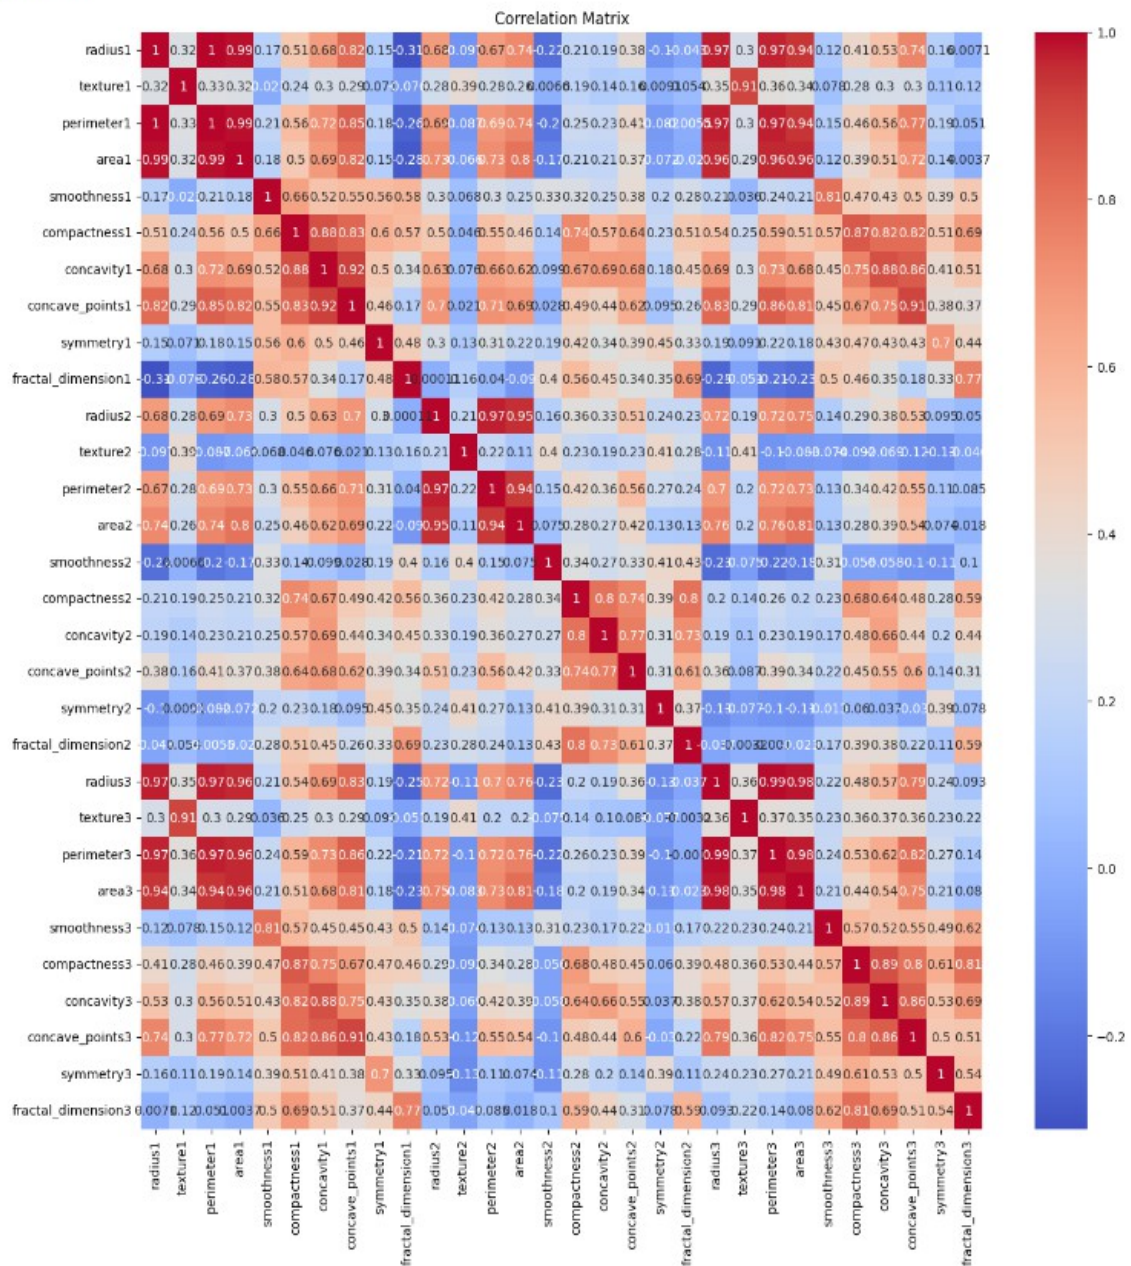
The T-tests determine how much of a significant difference there are between B & M. The Chi-square tests assess the independence between the features and how they relate to the target value.

Above we can see that the difference is very significant in the "radius1" and "perimeter1" fields, and less significant in other fields.



```
In [7]: # Create a Heatmap of the Correlation Matr

correlation_matrix = data.corr(numeric_only = True)
plt.figure(figsize=(15, 15))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```

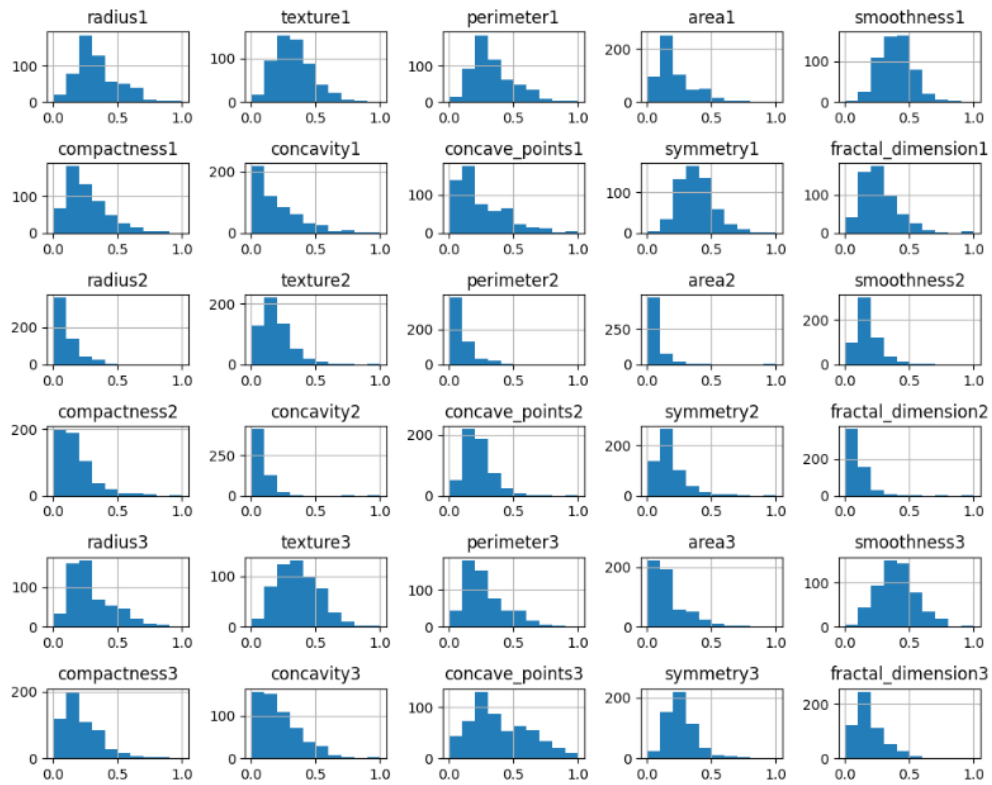


I also created a Heatmap that shows a correlation matrix. Here we can see how similar certain features are. Radius3 and Radius1 for example are very similar, while smoothness2 and fractal\_dimension3 are very different. Later in the feature selection stage we will see less values that are similar to each other appearing less.




In [6]: # Visualize the distribution of float values

```
data.drop('Diagnosis', axis=1).hist(figsize=(10, 8))
plt.tight_layout()
plt.show()
```



In [8]: # Generate box plots

```
data.drop('Diagnosis', axis=1).boxplot()
plt.title('Box Plots')
plt.xticks(rotation=90)
plt.show()
```



The box plot and distribution plot provide a visual representation of the distribution of data.

## Chapter 3 : Building ML algorithms

I will present the three basic learning algorithms I created below.

In all of my models I used a test size of 0.2 and a maximum iteration count of 1000.

When presented with the opportunity to set a random state, I consciously selected the number 42. By doing so, I aimed to achieve improved reproducibility of the algorithm, which is important for consistent results.

### Chapter 3.1 : Supervised model

For my first model, I chose to use a logistic regression approach. Logistic Regression is a commonly used and straightforward supervised learning algorithm, especially useful for binary classification tasks like my breast cancer detection algorithm. By training the logistic regression model with the labeled data I had, my aim was to find straightforward connections between the features and the target variable.

```
In [1]: # Supervised Learning: Logistic Regression
```

```
In [2]: # Imports
```

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
```

```
In [3]: data = pd.read_csv('wdbc.csv')
```

```
X = data.iloc[:, 1:]
y = data.iloc[:, 0]
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
```

```
y_pred = model.predict(X_test)
```

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
B	0.97	1.00	0.99	71
M	1.00	0.95	0.98	43
accuracy			0.98	114
macro avg	0.99	0.98	0.98	114
weighted avg	0.98	0.98	0.98	114

The supervised model performed exceptionally well, achieving an immediate impressive accuracy rating of 0.98. This outcome demonstrates the effectiveness and efficiency of the model in accurately predicting the target variable based on the labeled data.

## Chapter 3.2 : Unsupervised model

For my subsequent model, I opted for an unsupervised K-means clustering approach. Typically, this method is used when dealing with missing data or in scenarios where labeled information is unavailable. However, in this case, I chose to compare its outcomes to those of the first model.

While it is not be directly possible to directly compare the outputs of the unsupervised K-means clustering approach with the supervised model, I still can gain insights by combining these models with feature selection techniques and comparing their behaviour. By incorporating feature selection, we can assess the impact on performance and potentially gain a better understanding of the strengths and weaknesses of each model in relation to the dataset at hand.

Chapter 6 will provide an in-depth examination of feature selection techniques.

```
In [1]: # Unsupervised Learning: K-means Clustering

In [2]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans

data = pd.read_csv('wdbc.csv')

X = data.iloc[:, 1:]

inertia = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, n_init=10, random_state=42)
    kmeans.fit(X)
    inertia.append(kmeans.inertia_)


plt.plot(range(1, 11), inertia)
plt.title('Elbow Curve')
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.show()

kmeans = KMeans(n_clusters=3, n_init=10, random_state=42)
kmeans.fit(X)

data['Cluster'] = kmeans.labels_

print(data[['Diagnosis', 'Cluster']])
```

When we visualize the elbow graph we can see the Intertia decrease for each additinal cluster. This indicates that the clusters become more tightly packed, resulting in reduced distances between the data points and their respective cluster centers. Understanding the graph helps us create a balance between using meaningful structure within the data and avoiding unnecessary complexity.



	Diagnosis	Cluster
0	M	1
1	M	1
2	M	1
3	M	2
4	M	1
..	...	...
564	M	1
565	M	1
566	M	1
567	M	1
568	B	0

[569 rows x 2 columns]

The above picture of the elbow graph illustrates the impressive performance of the model.

## Chapter 3.3 : Semi-supervised model

For my third learning model, I decided to use a Semi-Supervised label propagation model.

I first loaded and prepared the data. The 'X' variable contains the features of the dataset (all columns except the first one), while 'y' contains the corresponding labels (the first column).

The 'train\_test\_split' function is used to split the data into labeled and unlabeled portions.

'X\_labeled' and 'y\_labeled' represent the labeled data, while 'X\_unlabeled' and 'y\_unlabeled' represent the unlabeled data. Here, 80% of the data is used for the unlabeled portion. I then created a LabelPropagation model. The model was then trained on the labeled data and produced the classification report below.

```
In [1]: # Semi-Supervised Learning : Label Propagation
```

```
In [2]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn import datasets
from sklearn.semi_supervised import LabelPropagation
from sklearn.metrics import classification_report

data = pd.read_csv('wdbc.csv')

X = data.iloc[:, 1:]
y = data.iloc[:, 0]

X_labeled, X_unlabeled, y_labeled, y_unlabeled = train_test_split(X, y, test_size=0.8, random_state=42)

model = LabelPropagation()
model.fit(X_labeled, y_labeled)

y_pred = model.predict(X_unlabeled)

print(classification_report(y_unlabeled, y_pred))
```

	precision	recall	f1-score	support
B	0.97	0.98	0.98	290
M	0.96	0.95	0.96	166
accuracy			0.97	456
macro avg	0.97	0.97	0.97	456
weighted avg	0.97	0.97	0.97	456

Here we can see a total accuracy of 0.97. The results are almost as good as in the first algorithm.

## Chapter 4 : Constructing and Selecting Features

After creating the three machine learning models, I explored what effect different feature selection methods had on these machine learning algorithms.

I choose five different feature selection algorithms : `SelectKBest_chi2`, `SelectKBest_f_classif`, `SelectKBest_mutual_info_classif`, `SelectFromModel_RandomForest` and `SelectFromModel_LinearSVC`. I then compared them with each of the models.

I then created five variations of each of the feature selection algorithms, with a different number of features to select : 5, 10, 15, 20 and 25.

adding the three models from chapter five (which do not use feature selection), this results in 78 ( $5 * 3 * 5 + 3$ ) differently trained models.

Doing this helped compare the models, the feature selection algorithms and their parameters.

To achive this, I used the original post-processed "wdbc.csv"-file and created the required 25 new CSV files uisng my "Compare Feature Selection" script below. These 25 new csv files would then be used by each model.

Saving the CSV files has proven to be highly advantageous in terms of optimizing both RAM and CPU usage. By loading the files individually as needed, I have successfully reduced the strain on the system's memory. Additionally, this approach eliminates the need to rerun feature selection algorithms multiple times with identical parameters, resulting in significant CPU savings.

```

In [33]: import pandas as pd
from sklearn.feature_selection import SelectKBest, chi2, f_classif, mutual_info_classif, SelectFromModel
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import LinearSVC

# LOAD CSV
data = pd.read_csv('wdbc.csv')

X = data.iloc[:, 1:]
y = data.iloc[:, 0]

# FEATURE COUNT
ks = [5, 10, 15, 20, 25]

for k in ks:

    # LIST OF METHODS
    methods = [
        ('SelectKBest_chi2', SelectKBest(chi2, k=k)),
        ('SelectKBest_f_classif', SelectKBest(f_classif, k=k)),
        ('SelectKBest_mutual_info_classif', SelectKBest(mutual_info_classif, k=k)),
        ('SelectFromModel_RandomForest', SelectFromModel(RandomForestClassifier(), max_features=k)),
        ('SelectFromModel_LinearSVC', SelectFromModel(LinearSVC(), max_features=k)),
    ]

    # APPLY METHOD
    for name, selector in methods:

        X_selected = selector.fit_transform(X, y)
        selected_indices = selector.get_support(indices=True)
        selected_features = X.columns[selected_indices]

        # VIEW SELECTED FEATURES

        print(f"Selected Features ({k}) for {name}:")
        fts = ""
        for feature in selected_features:
            if len(fts) > 0:
                fts += ", "
            fts += (feature)
        print(fts)
        print()

        selected_data = pd.DataFrame(X_selected, columns=selected_features)
        selected_data_with_target = pd.concat([selected_data, y], axis=1)

        # SAVE FILE

        path = f'wdbc_{name}_{k}.csv'
        selected_data_with_target.to_csv(path, index=False)

        df = pd.read_csv(path)
        last_column = df.iloc[:, -1]
        df = df.iloc[:, :-1]
        df.insert(0, 'Diagnosis', last_column)
        df.to_csv(path, index=False)

```

All created CSV Files can be viewed in the Github repo under the naming convention "wdbc\_{x}\_{y}.csv" where "x" represents the name of the feature selection algorithm and "y" represents the number of selected features.

"Methods" contains all the feature selection algorithms and "ks" contains all the amounts of features to be selected. This can be edited to create even more variations.

The features selected by each algorithm are shown in the "[Selected Features.txt](#)" file located in the coursework repository.



## Chapter 5 : Evaluating models and analysing the results

At this point I created the "All algorithms - full comparison with feature selection.ipynb" script to run the final above described test.

First I prepared the learning algorithms from Chapter 5 and made them easily accessible in my final evaluation algorithm below. Here I tested both the performance of the models with different feature selection settings.

```
In [1]: def check_accuracy_supervisedLearning(path):

    # ALGO A

    import pandas as pd
    from sklearn.model_selection import train_test_split
    from sklearn.linear_model import LogisticRegression
    from sklearn.metrics import classification_report
    from sklearn.metrics import precision_score

    data = pd.read_csv(path)

    X = data.iloc[:, 1:]
    y = data.iloc[:, 0]

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
    model = LogisticRegression(max_iter=1000)
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)

    precision = precision_score(y_test, y_pred, average='weighted')

    return precision
```

```
In [2]: def check_accuracy_unsupervisedLearning(path):

    # ALGO B

    import pandas as pd
    import matplotlib.pyplot as plt
    from sklearn.cluster import KMeans
    from sklearn.metrics import silhouette_score

    data = pd.read_csv(path)

    X = data.iloc[:, 1:]

    silhouette_scores = []

    for k in range(2, 11):
        kmeans = KMeans(n_clusters=k, n_init=10, random_state=42)
        kmeans.fit(X)
        labels = kmeans.labels_
        silhouette_scores.append(silhouette_score(X, labels))

    best_k = silhouette_scores.index(max(silhouette_scores)) + 2

    kmeans = KMeans(n_clusters=best_k, n_init=10, random_state=42)
    kmeans.fit(X)
    labels = kmeans.labels_

    precision = silhouette_score(X, labels)

    return precision
```

```
def check_accuracy_SemiSupervisedLearning(path):  
  
    # ALGO C  
  
    from sklearn.metrics import classification_report, precision_score  
    import pandas as pd  
    from sklearn.model_selection import train_test_split  
    from sklearn import datasets  
    from sklearn.semi_supervised import LabelPropagation  
    from sklearn.metrics import classification_report  
  
    data = pd.read_csv(path)  
  
    X = data.iloc[:, 1:]  
    y = data.iloc[:, 0]  
  
    X_labeled, X_unlabeled, y_labeled, y_unlabeled = train_test_split(X, y, test_size=0.2, random_state=42)  
  
    model = LabelPropagation()  
    model.fit(X_labeled, y_labeled)  
  
    y_pred = model.predict(X_unlabeled)  
  
    precision = precision_score(y_unlabeled, y_pred, average='macro')  
  
    return precision
```

I then ran the following code to check each algorithm, with each feature selection CSV and each their selected features :

In [5]:

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

FeatureSelectionPath = [
    "wdbc_SelectFromModel_LinearSVC_",
    "wdbc_SelectFromModel_RandomForest_",
    "wdbc_SelectKBest_chi2_",
    "wdbc_SelectKBest_f_classif_",
    "wdbc_SelectKBest_mutual_info_classif_"
]

FeatureSelectionLabels = [
    "SVC",
    "Forest",
    "KB-Chi2",
    "KB-Clsf",
    "KB-MI-Clsf",
]

FeatureSelectionCount = [5, 10, 15, 20, 25]

for i in range(0, 3):

    df = pd.DataFrame(0, index=FeatureSelectionCount, columns=FeatureSelectionLabels)

    indexlabel = 0

    for FeatureSelectionLabel in FeatureSelectionLabels:
        for FeatureSelectionCountCurrent in FeatureSelectionCount:

            path = FeatureSelectionPath[indexlabel] + str(FeatureSelectionCountCurrent) + "_".csv"

            val = 0;

            if i == 0:
                val = check_accuracy_supervisedLearning(path)
            if i == 1:
                val = check_accuracy_unsupervisedLearning(path)
            if i == 2:
                val = check_accuracy_SemiSupervisedLearning(path)

            df.loc[FeatureSelectionCountCurrent, FeatureSelectionLabel] = round(val, 4)

            indexlabel += 1

    fig, ax = plt.subplots(figsize=(8, 2))

    table = ax.table(cellText=df.values,
                    colLabels=df.columns,
                    rowLabels=df.index,
                    cellLoc='center',
                    loc='center',
                    cellColours=plt.cm.Greens(np.zeros_like(df.values))
                    )

    table.auto_set_font_size(False)
    table.set_fontsize(12)
    table.scale(1.2, 1.2)

    ax.axis('off')

    ax.set_title(["Supervised learning", "Unsupervised learning", "Semi-supervised learning"][i])
    plt.show()
```

Running the script produced the following tables :

#### Supervised learning

	SVC	Forest	KB-Chi2	KB-Clf	KB-MI-Clf
5	0.9569	0.9475	0.9475	0.9475	0.9475
10	0.9668	0.9569	0.9569	0.9569	0.9668
15	0.9652	0.9668	0.9569	0.9569	0.9569
20	0.9652	0.9668	0.9748	0.9748	0.9652
25	0.9652	0.9569	0.9748	0.9748	0.9748

#### Unsupervised learning

	SVC	Forest	KB-Chi2	KB-Clf	KB-MI-Clf
5	0.5014	0.5883	0.5825	0.5761	0.5883
10	0.4923	0.5742	0.5571	0.5571	0.5706
15	0.4402	0.5718	0.5293	0.5293	0.5293
20	0.4402	0.5718	0.4529	0.4529	0.4643
25	0.4402	0.5571	0.4186	0.4186	0.4186

#### Semi-supervised learning

	SVC	Forest	KB-Chi2	KB-Clf	KB-MI-Clf
5	0.9733	0.9538	0.9538	0.9538	0.9538
10	0.9671	0.9605	0.9605	0.9605	0.9605
15	0.9605	0.9605	0.9605	0.9605	0.9605
20	0.9605	0.9605	0.9554	0.9554	0.9673
25	0.9605	0.9605	0.9554	0.9554	0.9554

These tables display the output of each of the learning algorithms, with both the feature selection ( X-Axis) and the number of selected features (Y-Axis).

In the supervised and semi-supervised tables I displayed the precision scores, while the unsupervised learning table is based on the silhouette score.

viewing the tables and comparing them with the outputs given in chapter 5 (the outputs without any feature selection), we can clearly see that the supervised learning model, without any feature selection works best:

	precision	recall	f1-score	support
B	0.97	1.00	0.99	71
M	1.00	0.95	0.98	43
accuracy			0.98	114
macro avg	0.99	0.98	0.98	114
weighted avg	0.98	0.98	0.98	114

This was expected because the unsupervised and semi-supervised models tend to work better than the supervised model in cases where there is a lack of data (for example: lack of labels or noise). In this case, the dataset I worked with had plenty of data, good features, and was well-prepared.

When looking at the feature selection, it's interesting to see how well the model performs with just five features instead of using the full 30 features.

In the case of the SVC feature selection, the supervised model continues to have an impressive accuracy of 0.95. This minor decrease in accuracy compared to the full model demonstrates the efficiency of the feature selection process. Despite a significant reduction in the number of features, the model maintains a high level of performance, displaying the importance of selecting the most influential features for optimal results.

In larger datasets aiming to address this problem, using feature selection will reduce complexity and training time. SVC works well when using a small number of features, but if more than 15 features are used, any of the KB-\* feature selection systems work better.

Unfortunately, the RandomForest classifier doesn't perform well enough with a small number of features to beat SVC, and it doesn't work well enough with a large number of features to beat any of the KB-\* algorithms. It even seems to get worse after 20 features, both in supervised and unsupervised learning.

## Chapter 6 : Conclusion

The course and this task allowed me to understand the power that is machine learning. It is incredible to see how fast algorithms can create a deep understanding of raw data, and apply it to solve problems like these.

Im certain machine learning will become one of the most prevalent tools in medicine. The potential of this technology to save lives is incredible.

## Chapter 7 : Sources

1. [https://www.krebsdaten.de/Krebs/EN/Content/Cancer\\_sites/Breast\\_cancer/breast\\_cancer\\_node.html](https://www.krebsdaten.de/Krebs/EN/Content/Cancer_sites/Breast_cancer/breast_cancer_node.html)
2. <https://archive.ics.uci.edu/dataset/17/breast+cancer+wisconsin+diagnostic>
3. <https://github.com/University-of-London/csm010-aml-coursework-KaiBowers99>